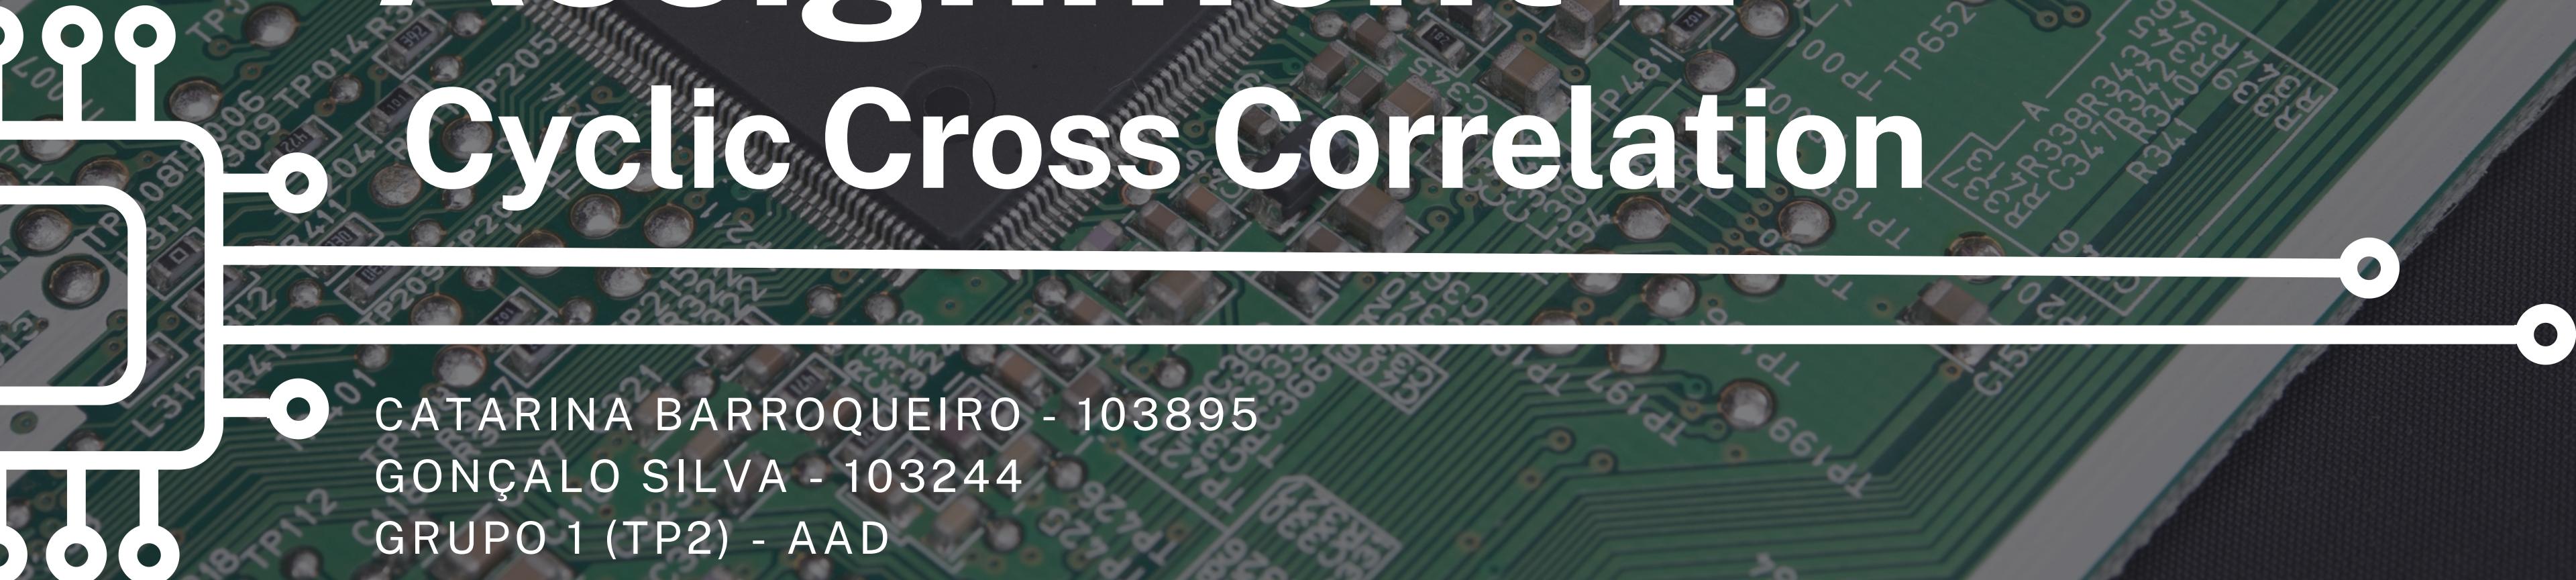


Assignment 2

Cyclic Cross Correlation



CATARINA BARROQUEIRO - 103895
GONÇALO SILVA - 103244
GRUPO 1 (TP2) - AAD

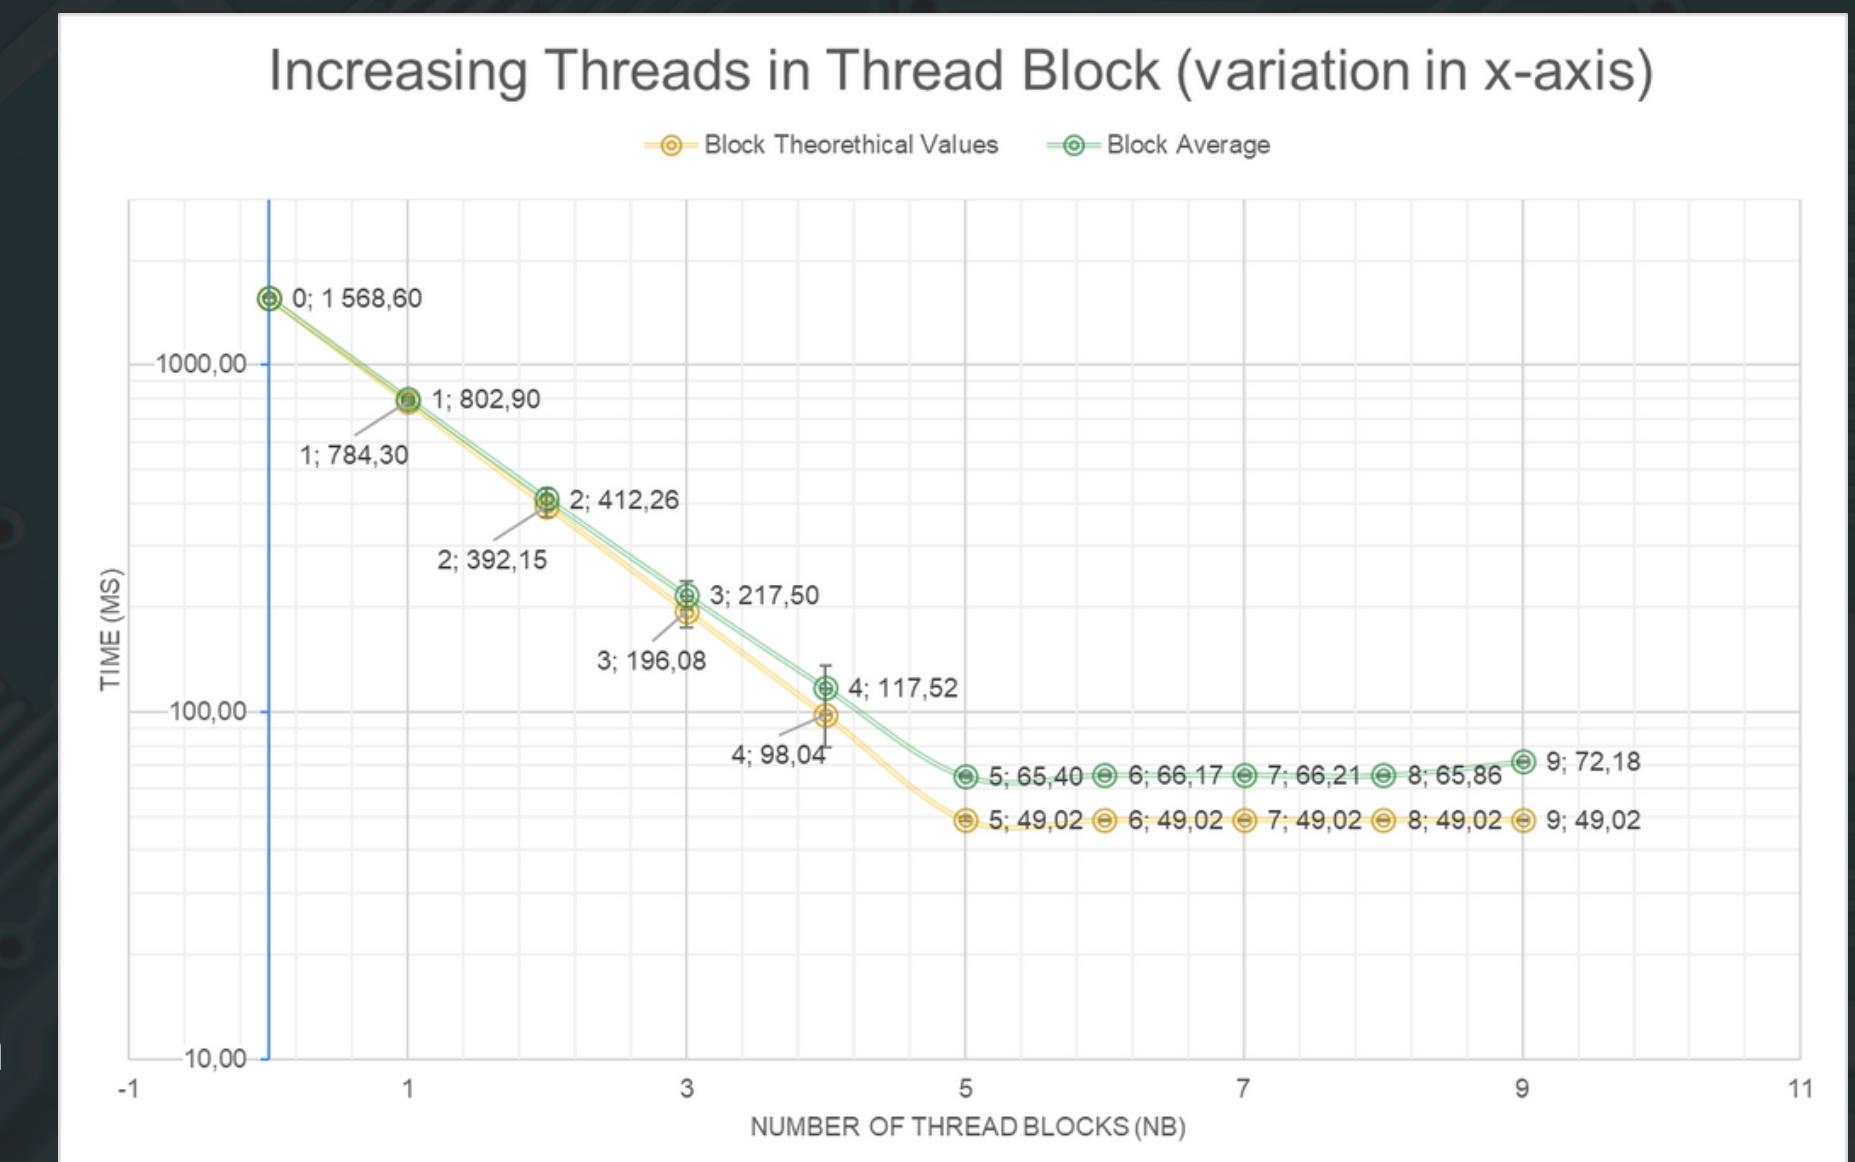
GOAL & OPTIMAL NUMBER OF THREADS PER BLOCK

GOAL

- Optimize circular cross-correlation thread launch grid, investigate GPU performance factors, and draw conclusions on computation efficiency when offloading to a GPU.

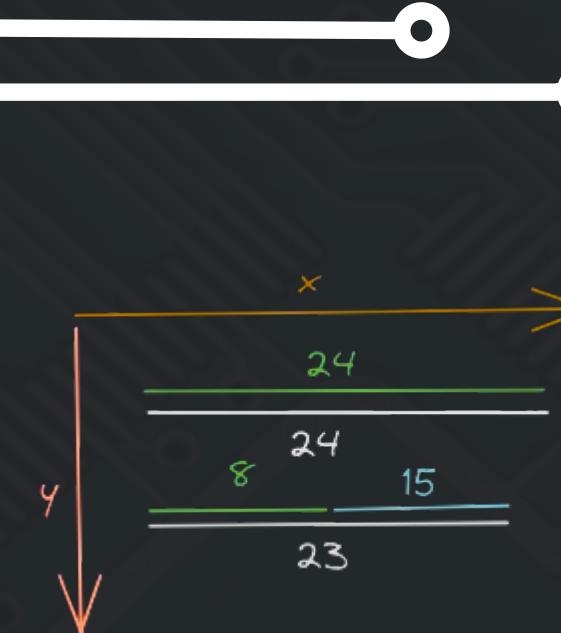
OPTIMAL NUMBER OF THREADS PER BLOCK

- Obtained the optimal configuration with grid dimension 2048, block dimension 32, yielding an average time of 65.45ms with a standard deviation of 0.663ms.



PROCESS

- Modifying Block dimensions to an **x-y** field generated huge differences in the average time of executions
- A cause can be the execution of each thread in a thread block, being processed linearly, even though appearing in a matrix format, it will always be executed by the SM in a 1D order
- Variance was also heavily impacted



- Another cause can be the execution order of each thread being on a first come first serve, which may be leading threads to always be in load while starving others
- Finally, with threads divided in the **x-y** axis, will lead the data of the program to be more disperse, causing more cache misses, thus slowing the program

Messing with X-Y Threads per Block													
GridDimX	GridDimY	GridDimZ	BlockDimX	BlockDimY	BlockDimZ	Run 1 (ms)	Run 2 (ms)	Run 3 (ms)	Run 4 (ms)	Run 5 (ms)	Average	Variance	
2048	1	1	32	1	1	65,12	66,06	64,54	66,12	65,43	65,45	0,663	
2048	1	1	16	2	1	60,86	68,79	68,80	68,80	68,81	67,21	3,5509	
2048	1	1	8	4	1	51,91	69,09	69,12	69,10	69,11	65,67	7,6898	
2048	1	1	4	8	1	83,49	83,41	63,18	62,88	83,34	75,26	11,165	
2048	1	1	2	16	1	144,10	108,80	143,20	142,70	144,20	136,60	15,5533	
2048	1	1	1	32	1	143,60	143,90	144,00	144,50	142,80	143,76	0,6269	
Launching Number of Thread Blocks (nB)			Launching Thread per Block (nT)										

PROCESS

- Modifying Grid dimensions to an **x-y** field generated small differences. This may be due to the execution being handled by the scheduller, which is more consistent
- Another factor is the grid only determining the number of thread blocks to be scheduled, not having impact in the handling of the threads and execution inside the block

Messing with X-Y Block Grid												
GridDimX	GridDimY	GridDimZ	BlockDimX	BlockDimY	BlockDimZ	Run 1 (ms)	Run 2 (ms)	Run 3 (ms)	Run 4 (ms)	Run 5 (ms)	Average	Variance
2048	1	1	32	1	1	65,12	66,06	64,54	66,12	65,43	65,45	0,663
1024	2	1	32	1	1	68,78	68,80	68,82	68,83	68,80	68,81	0,0195
512	4	1	32	1	1	68,81	68,81	68,80	68,83	68,80	68,81	0,0122
256	8	1	32	1	1	68,78	68,79	68,84	68,81	66,83	68,41	0,8835
128	16	1	32	1	1	68,82	68,80	68,78	68,79	68,81	68,80	0,0158
64	32	1	32	1	1	68,80	68,81	68,83	68,79	68,82	68,81	0,0158
32	64	1	32	1	1	68,78	68,80	68,79	68,82	68,82	68,80	0,0179
16	128	1	32	1	1	68,82	68,82	68,83	68,81	68,79	68,81	0,0152
8	256	1	32	1	1	68,79	68,78	68,80	68,81	68,79	68,79	0,0114
4	512	1	32	1	1	68,78	68,76	68,80	68,76	68,84	68,79	0,0335
2	1024	1	32	1	1	68,81	68,78	68,83	68,83	68,81	68,81	0,0205
1	2048	1	32	1	1	68,79	68,80	68,79	68,82	68,80	68,80	0,0122
Launching Number of Thread Blocks (nB)			Launching Thread per Block (nT)									

GPU vs CPU

- **CPU EXECUTION:**
 - **Average Time :** 3103 ± 1.8257 ms
- **GPU:**
 - **Average Time :** 200.08 ± 3.754 ms
 - **Memory Host -> Device:** 134.50 ± 3.6968 ms
 - **Memory Device -> Host :** 0.130 ± 0.0034 ms
 - **Processing Time :** 65.45 ± 0.663 ms
- **SPEED-UP:**
 - **Upper bound:** 15.566 times
 - **Lower bound:** 15.533 times

Executing the program in the GPU is around 16 times faster than on the CPU

CPU Execution						
Run 1 (ms)	Run 2 (ms)	Run 3 (ms)	Run 4 (ms)	Run 5 (ms)	Average	Variance
3120	3101,00	3102,00	3104,00	3105,00	3103,00	1,826

Memory Transfer							
Direction	Run 1 (ms)	Run 2 (ms)	Run 3 (ms)	Run 4 (ms)	Run 5 (ms)	Average	Variance
Host -> Device	136	138,00	133,00	130,00	137,00	134,50	3,6968
Device -> Host	0,127	0,129	0,135	0,128	0,128	0,130	0,0034